

# **Deductive Databases**

## **Semantic Optimization of Queries**

# Outline

---

- ▶ Semantic Query Optimization (SQO)
  - ▶ A quick refresher of basic notions
- ▶ Two-Phased Approach to SQO
  - ▶ Semantic compilation
  - ▶ Example scenarios
- ▶ Partial Subsumption
- ▶ Extended forms
- ▶ Semantically Constrained Axioms

# Semantic Query Optimization

---

- ▶ Use semantic knowledge (e.g. integrity constraints) for transforming a query in a more efficient form than the original version
- ▶ The strategy chosen for optimization and its cost are important considerations
  - ▶ We will study a two-phased approach to semantic query optimization

# A quick refresher of basic notions

---

- ▶ A **literal** is either an atomic formula or the negation of an atomic formula

- ▶ A **clause** is a disjunction of literals which has the form

$$S_1 \vee \dots \vee S_m \vee \neg R_1 \vee \dots \vee \neg R_n$$

- ▶ where each  $S_i$  and  $R_j$  are atomic formulas
- ▶ Clauses can be written in an equivalent form using implication
  - ▶  $S_1, \dots, S_m \leftarrow R_1, \dots, R_n$
  - ▶ All the variables in a clause are assumed to be universally quantified

# A quick refresher of basic notions

---

- ▶ In logic programming, a program consists of clauses
  - ▶ Prolog restricts clauses to be **Horn**
  - ▶ They are adequate for specifying a large class of database applications
- ▶ A clause is **Horn** if  $m \leq 1$  and disjunctive otherwise
$$S \leftarrow R_1, \dots, R_n$$
  - ▶ where each  $S$  and  $R_j$  are atomic formulas
- ▶ A Horn clause is one whose head consists of at most a single atom

# A quick refresher of basic notions

---

- ▶ A **Horn** clause is one whose head consists of at most a single atom
  - ▶ A **goal** clause has a null head;
  - ▶ The **null** clause has both a null body and a null head;
    - ▶ It represents a contradiction
  - ▶ A clause is **definite** if the head of the clause contains exactly one atom;
  - ▶ A **unit** clause is a definite clause with a null body;
  - ▶ A **ground unit** clause is a unit clause all of whose arguments are constants.

# A quick refresher of basic notions

---

- ▶ We use operators to denote relations corresponding to widely used arithmetic operations
  - ▶ such as  $>$ ,  $<$ ,  $=$ ,  $\neq$ ,  $\leq$ ,  $\geq$
  - ▶ These relations are **termed evaluable** (or built-in) relations in contrast to (**nonevaluable**) relations defined as part of the database
- ▶ A clause is **range-restricted** if every variable in the head also appears in the body as arguments of nonevaluable relations
- ▶ A clause is called **recursive** if the same relation symbol appears both in the head and the body

# A quick refresher of basic notions

---

- ▶ The notion of deductive database is based on the proof-theoretic approach to databases
- ▶ We limit our discussion to the basic components of a deductive databases
  - ▶ Facts
  - ▶ Deductive Rules (Axioms)
  - ▶ Integrity Constraints



# Facts

---

- ▶ Database facts are expressed as ground unit clauses
- ▶ The fact that a tuple  $\langle a_1, \dots, a_n \rangle$  is a member of the relation  $R$  is expressed by the atomic literal

$$R(a_1, \dots, a_n) \leftarrow$$

- ▶ Such facts comprise the extensional database, **EDB**, and such a relation  $R$  is called an extensional (stored) relation

# Deductive Rules

---

- ▶ The intensional database, **IDB**, contains deductive rules
  - ▶ Relations defined by deductive rules are called intensional relations

## Horn clause

Emp-Dept(ename,dname)

← Emp(eid,ename,did), Dept(did,dname)

- ▶ This is a rule defining the (intensional) Emp-Dept relation in terms of the (extensional) Emp and Dept relations
- ▶ In general, a rule

$$S \leftarrow R_1, \dots, R_n$$

- ▶ defines a relation  $S$  in terms of the relations  $R_1, \dots, R_n$

# Integrity Constraints

---

- ▶ Integrity constraints contain rules expressing restrictions that the database must satisfy
  - ▶ A large class of integrity constraints are expressible as Horn clauses
- ▶ The following are examples on the extensional relation `Supplier(Sno,Sname,Item)`

## Integrity Constraints

$y1=y2 \leftarrow \text{Supplier}(x, y1, z1), \text{Supplier}(x, y2, z2)$

$\leftarrow \text{Supplier}(x, y, \text{'gun'}), \text{Supplier}(x, z, \text{'butter'})$

# Integrity Constraints: Examples

---

$y1=y2 \leftarrow \text{Supplier}(x, y1, z1), \text{Supplier}(x, y2, z2)$

- ▶ It represents the functional dependency of supplier name on supplier number

$\leftarrow \text{Supplier}(x, y, \text{'gun'}), \text{Supplier}(x, z, \text{'butter'})$

- ▶ It states that no supplier supplies both the object 'gun' and the object 'butter'

# Integrity Constraints

---

- ▶ Integrity constraints play an important role in checking update validity
- ▶ They are not needed in answering queries over definite deductive databases
- ▶ Semantic query optimization builds upon the premise that integrity constraints are useful for query evaluation
  - ▶ They can be used in transforming a query into a form that is more efficient to evaluate

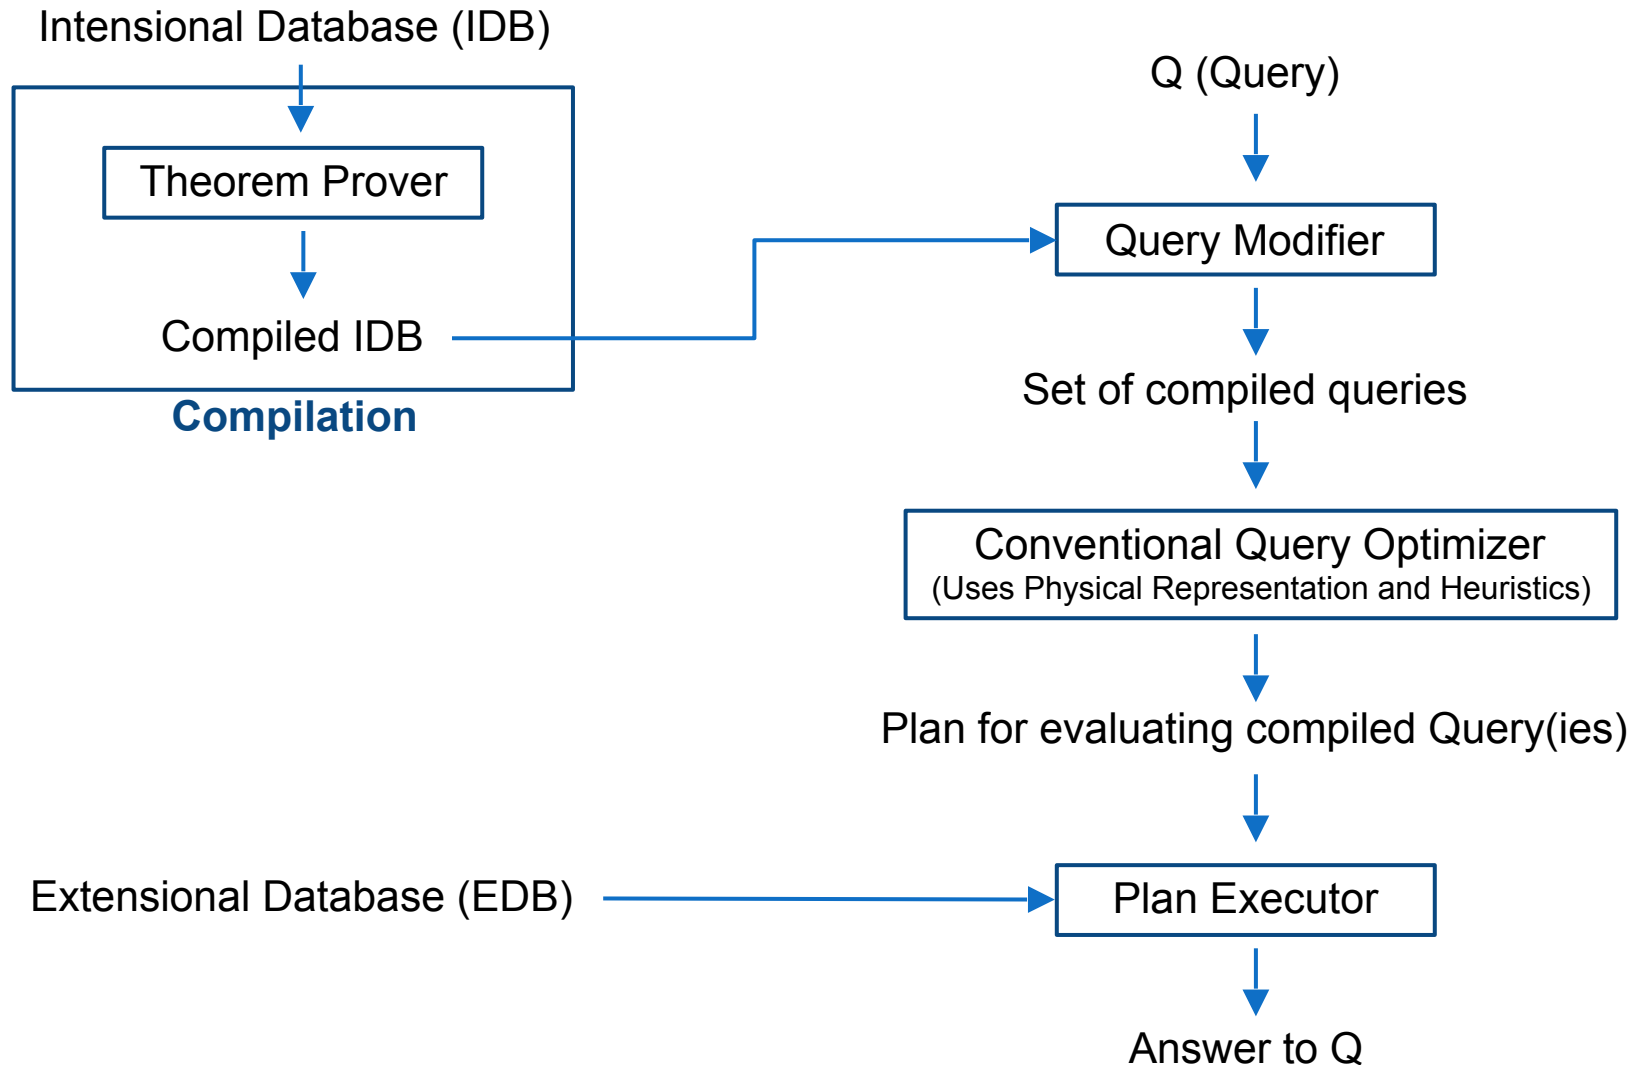
# Compiled Approach

---

- ▶ In (non-recursive) deductive databases, query evaluation can be partitioned into
  - ▶ a compilation phase followed by a modification and an evaluation phase
- ▶ With the compiled approach a query is evaluated by
  - ▶ first modifying the query using the compiled (intensional) database, and then
  - ▶ optimizing the result using conventional techniques

# Compiled Approach

---



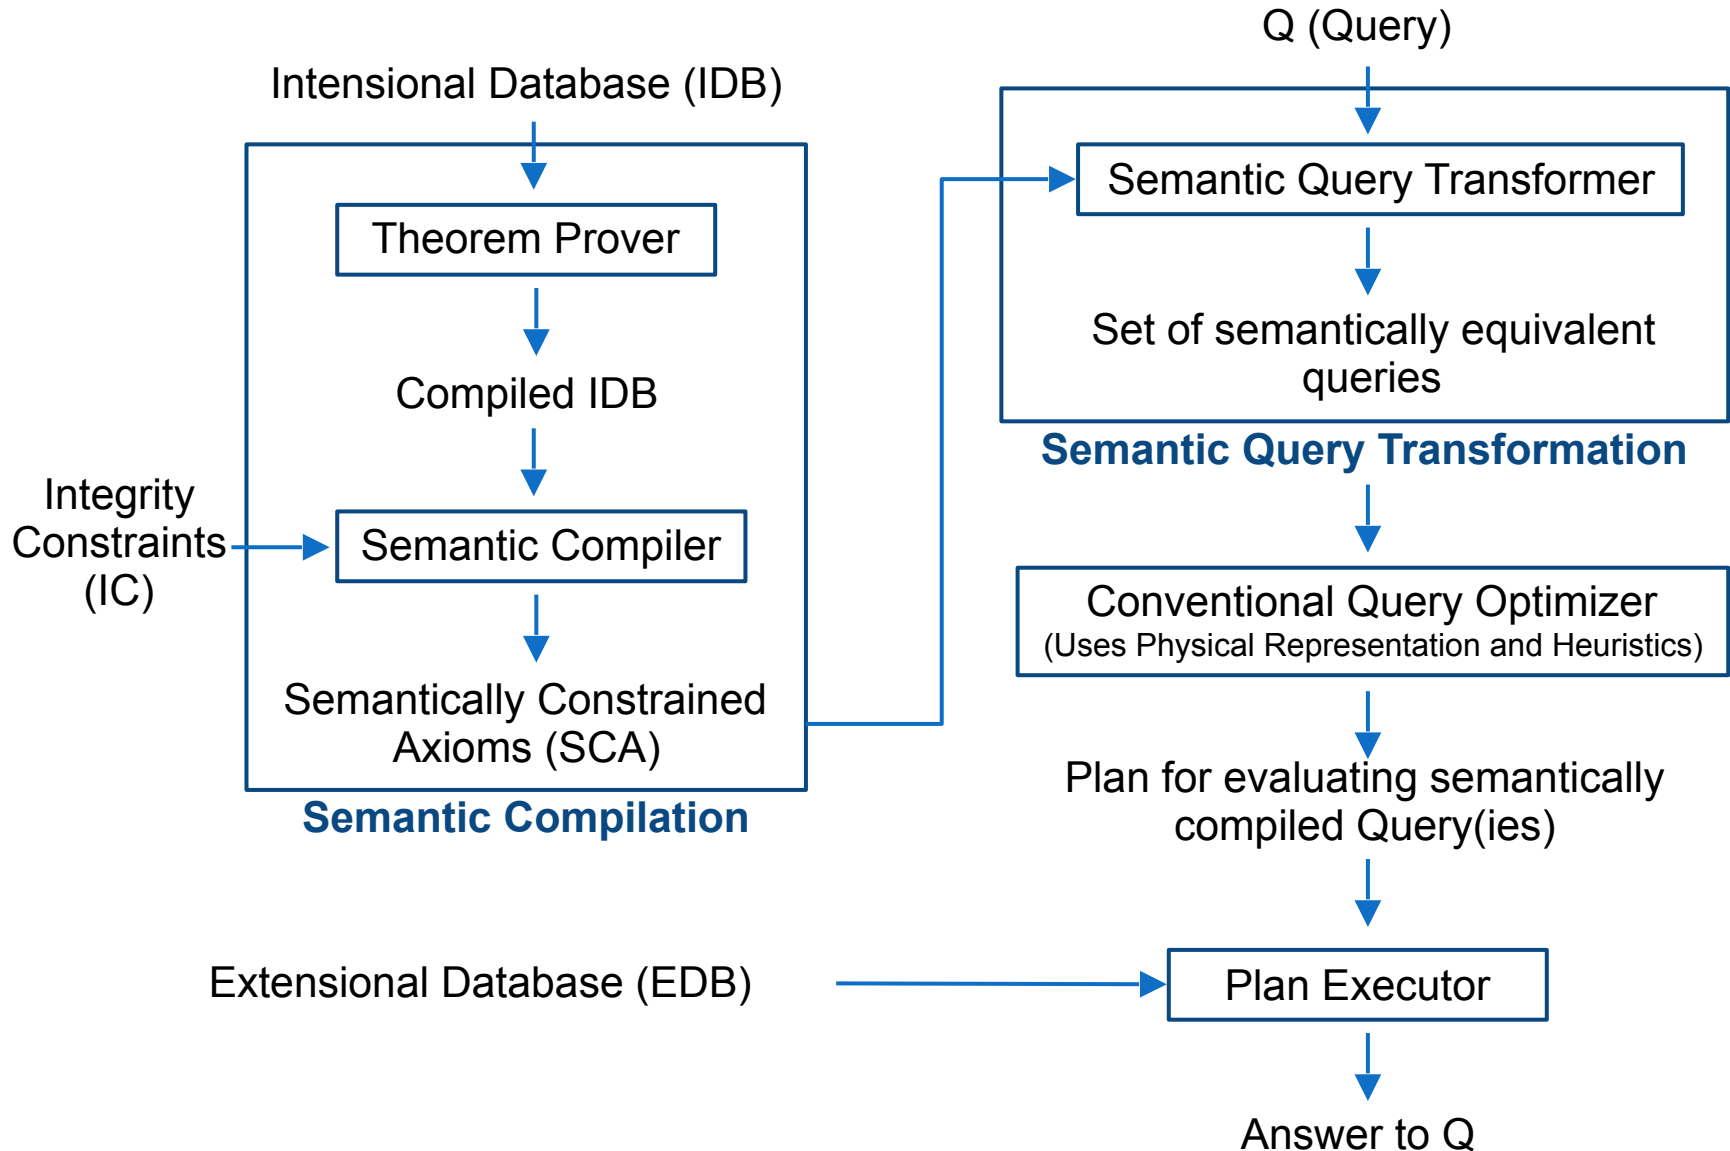
# Two-Phased Approach to SQO

---

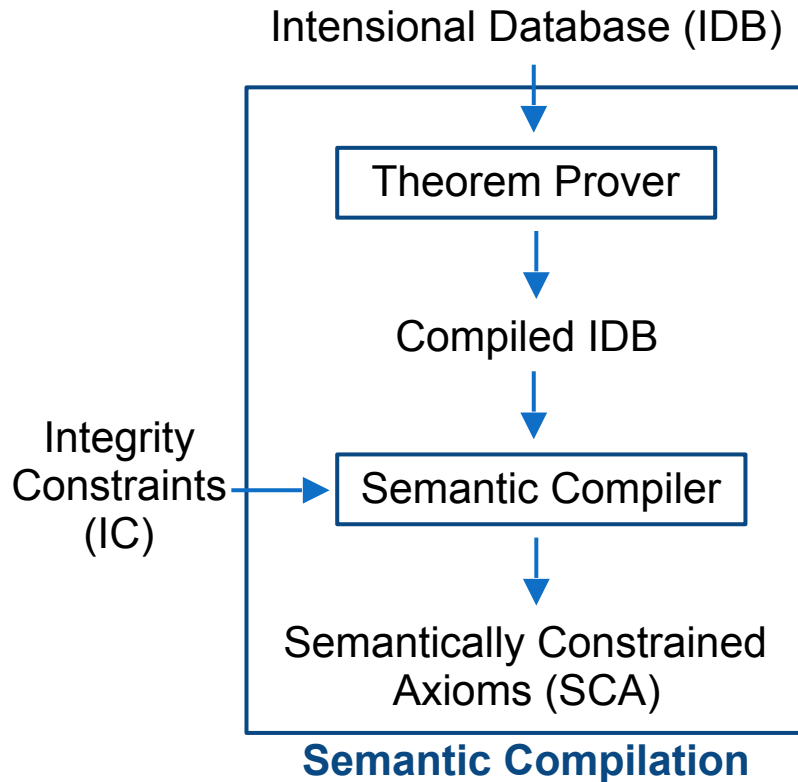
- ▶ The Two-Phased approach
  - ▶ Uses integrity constraints for optimizing a query over a deductive database
  - ▶ Generates several semantically equivalent queries as a result of the query modification stage, instead of a single query
- ▶ This entails changes to both
  - ▶ the compilation stage (now, semantic compilation phase)
  - ▶ the query modification stage (now, semantic transformation phase)



# Two-Phased Approach to SQO

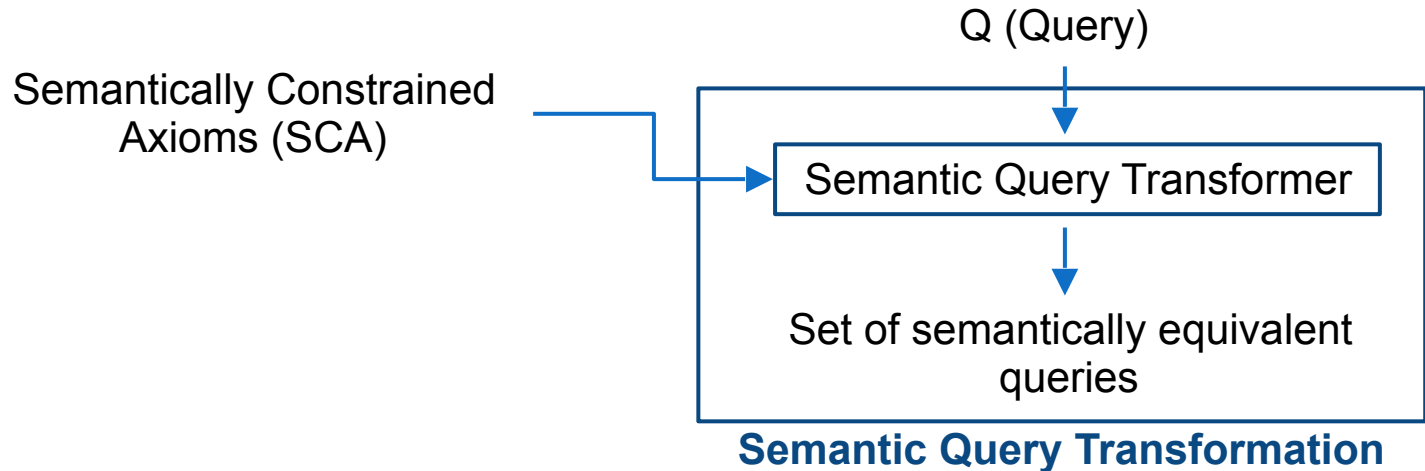


# Two-Phased Approach to SQO



- ▶ During the semantic compilation phase, integrity constraint fragments, called **residues**, are computed and associated with deductive rules
- ▶ The result is a set of **semantically constrained axioms (SCA)** which are filtered and stored for later use

# Two-Phased Approach to SQO



- ▶ When a query is given to the system, the semantic transformer uses these stored residues to generate semantically equivalent queries that *may be processed faster than the original query*

# Types of Transformations

---

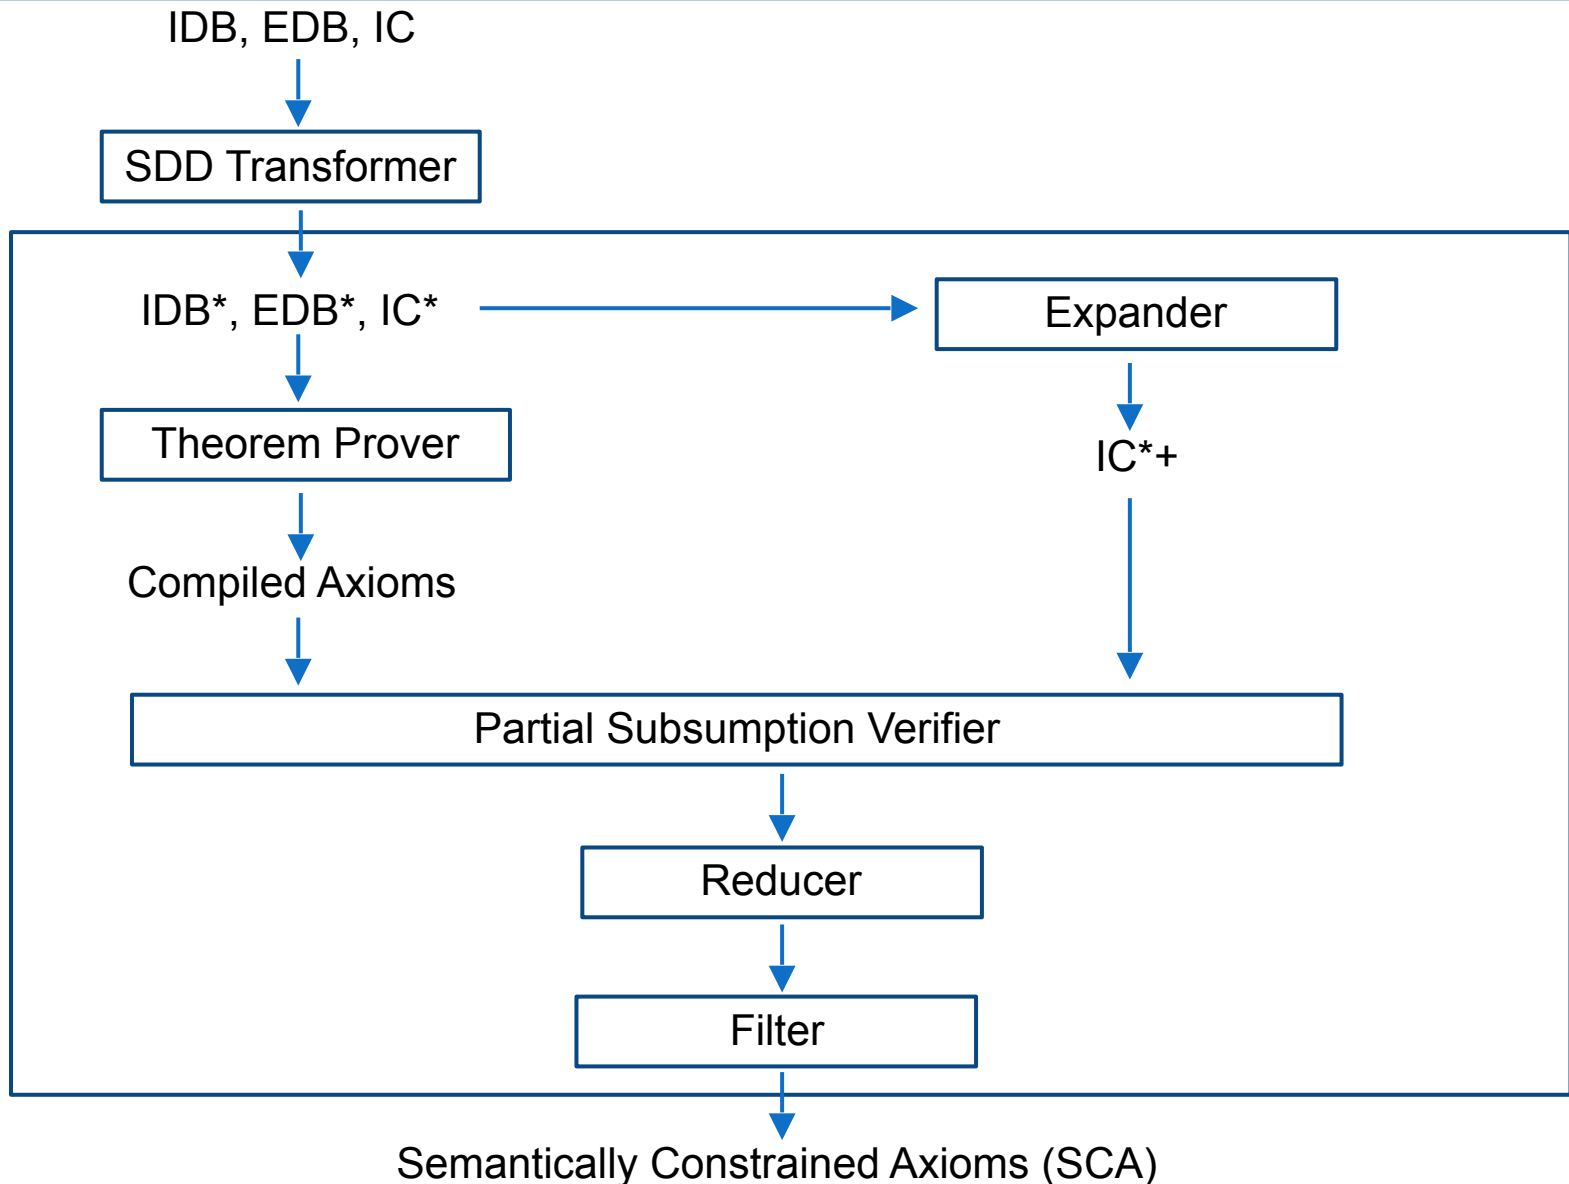
- ▶ Transformations to queries are performed using various heuristics
  - ▶ Literal elimination
    - ▶ hence a join, which is an expensive relational operation
  - ▶ Restriction introduction
    - ▶ may make it possible to use a range search, thereby reducing the cost of evaluation
  - ▶ Literal introduction
    - ▶ in some cases the introduction of a small instantiated relation for an additional join can reduce the amount of computation
  - ▶ Transformation that answers the query
  - ▶ Detection of unsatisfiable conditions

# Semantic Compilation

---

- ▶ Semantic compilation can be described informally as the process of retaining relevant (or processed) fragments of integrity constraints
  - ▶ A step in which all useful information has been extracted from the integrity constraints
- ▶ Semantic compilation is performed only once for a given deductive (or relational) database
  - ▶ Is used over relatively stable database components
- ▶ The primary objective of semantic compilation is to associate integrity constraint fragments, called residues, with compiled axioms

# Semantic Compilation Architecture



# An example scenario

---

**EDB:** Ships(sname,owner,stype,draft,deadweight,capacity, registry)

Owners(oname,location,assets,business)

**IC:**

IC1:  $\leftarrow$  Owners(x1,'iceland',x3,x4)

IC2:  $x2 = \text{'onassis'} \leftarrow$  Ships(x1,x2, 'supertanker',x4,x5,x6,x7)

**Q:**

Q1:  $\leftarrow$  Owners(x1\*,'iceland',x3,x4\*), (x3>1000000)

Q2:  $\leftarrow$  Ships(x1,x2,'superthanker',x4,x5,x6,x7\*)

Q3:  $\leftarrow$  Owners(y1\*,y2,y3,y4), Ships(x1,y1\*, 'superthanker',x4,x5,x6,x7)

# An example scenario

EDB: Ships(sname,owner,stype,draft,deadweight,capacity, registry)

Owners(oname,location,assets,business)

IC1:  $\leftarrow$  Owners(x1,'iceland',x3,x4)

□ «There are no owners with 'iceland' as their business location»

IC2:  $x2 = \text{'onassis'} \leftarrow$  Ships(x1,x2, 'supertanker',x4,x5,x6,x7)

□ «All 'supertankers' are owned by 'onassis'»

- ▶ We show now how these IC may be useful in query processing
  - ▶ IC1 is useful only if Owners appears in the query, and
  - ▶ IC2 is useful only if Ships appears in the query



# An example scenario

---

- ▶ Let now suppose the following query

Q1:  $\leftarrow \text{Owners}(x1^*, 'iceland', x3, x4^*), (x3 > 1000000)$

- It asks for the names and businesses of owners who are located in 'Iceland' and whose assets are greater than 1 million

- ▶ A relational database system may solve such a query by a table lookup of the Owners relation or by using an index on the Owners relation
- ▶ However, it follows from IC1 that Q1 cannot have any answers, so no search is needed at all

# An example scenario

---

- ▶ Let now suppose the following query

Q2:  $\leftarrow \text{Ships}(x1, x2, \text{'superthanker'}, x4, x5, x6, x7^*)$

- It asks for the registry information of all 'supertankers'
- ▶ Assuming that the relation is not indexed on stype, a table lookup is necessary
- ▶ However, by using IC2 it is sufficient to look for x2 values of 'onassis' and the corresponding registry information
  - ▶ IC2 introduces a selection criterion for the owner attribute

# An example scenario

---

- ▶ Let now suppose the following query

Q3:  $\leftarrow \text{Owners}(y1^*, y2, y3, y4), \text{Ships}(x1, y1^*, \text{'superthanker'}, x4, x5, x6, x7)$

- It asks for the owner names of all ships of the type 'supertanker'

- ▶ Using IC2 one can establish the value of  $y1^*$  as 'onassis'.
  - ▶ However, it is necessary to confirm the presence of such a tuple in both the relations Ships and Owner
- ▶ An index is introduced on a join attribute, which also happens to be the output attribute in this example
  - ▶ In the presence of appropriate existential and inclusion integrity constraints, both the join and the lookup can be eliminated

# Another example scenario

---

EDB: Emp(ssno,salary,deptno,age)

Dept(deptno,manager,floor)

Sales(deptno,item,vol)

IDB: Highsales(x1,x2,x3,y2,y3)  $\leftarrow$   
Dept(x1,x2,x3), Sales(x1,y2,y3), y3 > 100000

HMgrProf(x2,y2,y4)  $\leftarrow$  Emp(x2,y2,y3,y4), Highsales(x1,x2,x3,x4,x5)

IC:

IC1:  $\leftarrow$  Dept(x,y,2)

IC2: (y > 40000)  $\leftarrow$  Emp(x,y,z,u), (u > 50)

Q:

Q1:  $\leftarrow$  Highsales(x\*,y,2,z\*,u)

# Another example scenario

---

- ▶ A standard method for dealing with such a database is to reduce all definitions of intensional relations to extensional relations
- ▶ As long as the axioms (rules) in IDB are not recursive, the body of every axiom can be reduced to purely extensional relations using the **compiled method**

```
HMgrProf(x1,x2,x4) ←  
    Emp(x1,x2,x3,x4),Dept(y1,x1,y2),Sales(y1,z2,z3),z3>100000
```

# Another example scenario

---

- ▶ Let now suppose the following query

Q1:  $\leftarrow \text{Highsales}(x^*, y, 2, z^*, u)$

- It asks for the deptno and item values that are sold in large quantity (i.e., volume > 100000) by that department

- ▶ A deductive database system may solve this query by transforming it using the rule for Highsales to

Q1':  $\leftarrow \text{Dept}(x^*, y, 2), \text{Sales}(x^*, z^*, u), (u > 100000)$

- ▶ It follows from IC1 and Q1' that there are no answers

# Subsumption

---

- ▶ A substitution ( $\sigma$ ) is a finite set  $\{t_1/v_1, \dots, t_n/v_n\}$ , where  $v_i$ 's are unique variables and  $t_i$ 's are terms
- ▶  $C\sigma$  is the clause obtained by replacing each occurrence of the variable  $v_i$  in  $C$  which is also in  $\sigma$  by the term  $t_i$
- ▶ A clause  $C$  is a subclause of  $D$  if every literal in the clause  $C$  is also in  $D$

*A clause  $C$  **subsumes** a clause  $D$  if there is a substitution  $\sigma$  such that  $C\sigma$  is a subclause of  $D$*

# Subsumption

---

## ▶ Example 1

$R(x,b) \leftarrow P(x,y), Q(y,z,b)$

$R(a,b) \leftarrow P(a,z), Q(z,z,b), S(a)$

$\sigma = \{a/x, z/y\},$


## ▶ Example 2

IC1:  $\leftarrow \text{Owners}(x1, \text{'iceland'}, x3, x4)$

Q1:  $\leftarrow \text{Owners}(x1^*, \text{'iceland'}, x3, x4^*), (x3 > 1000000)$



# Refusion Tree



$R(x,b) \leftarrow P(x,y), Q(y,z,b)$

$R(a,b) \leftarrow P(a,k1), Q(k1,k1,b), S(a)$

# The Role of Partial Subsumption

---

- ▶ The semantic compilation phase occurs before any queries are posed
  - ▶ we must use the integrity constraints and the axioms for the relations to anticipate later modifications to queries
- ▶ The essence of **partial subsumption** is to apply the subsumption algorithm to an integrity constraint and the body of an axiom
  - ▶ a subclause of the integrity constraint might subsume the body of the axiom
  - ▶ Instead of the null clause, a fragment of the integrity constraint remains at the bottom of a refutation tree
    - ▶ Such a fragment is called a residue


# The Role of Partial Subsumption

IC1:  $\leftarrow \text{Owners}(x1, \text{'iceland'}, x3, x4)$  (C)

Axiom:  $\leftarrow \text{Owners}(y1, y2, y3, y4)$  (D)

$\text{Owners}(x1, \text{'iceland'}, x3, x4)$

$\text{Owners}(k1, k2, k3, k4) \leftarrow$



- ▶ A variable is needed in place of the constant “iceland”
  - ▶ Consider C into the expanded form, denoted by  $C^+$ , where C and  $C^+$  are logically equivalent

*An IC **partially subsumes** an axiom A if IC does not subsume the body of A, but a subclause of  $IC^+$  subsumes the body of A*

# Expanded Forms

---

- ▶ An integrity constraint is expanded by substituting variables in place of constants systematically, in order to apply subsumption and resolve as many literals as possible

## Steps of expansions

- ▶ The expanded clause  $C^+$  is obtained from the clause  $C$  by modifying the body of  $C$  as follows:
  1. An evaluable relation (predicate) that contains a constant and a variable is modified to two relations with the original relation containing two variables
    - ▶  $(u > c)$  is replaced by  $(u > x_1), (x_1 \geq c)$

# Expanded Forms

---

## Steps of expansions

- ▶ The expanded clause  $C^+$  is obtained from the clause  $C$  by modifying the body of  $C$  as follows:
  2. An extensional relation that contains a constant or a variable that has occurred previously (that is, to its left in the clause) is modified by changing the constant or the variable to a new variable and adding an equality consisting of the constant or the variable and the new variable introduced
- ▶  $IC1^+ : \leftarrow \text{Owners}(y1, x1, y3, y4), (x1 = \text{'iceland'})$

# Expanded Forms

IC1+:  $\leftarrow \text{Owners}(y1, x1, y3, y4), (x1 = \text{'iceland'})$  (C)

Axiom:  $\leftarrow \text{Owners}(y1, y2, y3, y4)$  (D)

$\leftarrow \text{Owners}(y1, x1, y3, y4), (x1 = \text{'iceland'})$

$\text{Owners}(k1, k2, k3, k4) \leftarrow$

$\{k1/y1, k2/x1, k3/y3, k4/y4\}$


$\leftarrow k2 = \text{'iceland'}$

- ▶ We can interpret this clause as “y2 cannot be iceland for  $\text{Owners}(y1, y2, y3, y4)$ ”

# Samantically Constrained Axiom

- ▶ Given an IC and an axiom A, apply the subsumption to IC+ and the body of A until no more resolutions are possible, and let B be the clause at the bottom of a refutation tree

- ▶ Then  $(B-)\sigma^{-1}$  that is, the clause obtained by the back substitution of the reduction of B is a **residue** of IC and A



$\text{Owners}(y1,y2,y3,y4) \leftarrow \text{Owners}(y1,y2,y3,y4), \{\leftarrow y2='iceland'\}$

# Meaning of SCA

---

- ▶ A compiled intensional axiom  $H \leftarrow P_1, \dots, P_m$  (terms omitted) is transformed to its semantically constrained form  $H \leftarrow P_1, \dots, P_m \{R_1, \dots, R_m\}$ , where the  $R_i$  are residues
- ▶ Let  $R_i$  have the form  $G \leftarrow F_1, \dots, F_k$  and define  $R_i'$  as  $\text{not}(F_1, \dots, F_k, \text{not}(G))$ 
  - ▶ In the case of a unit residue,  $R_i'$  is  $G$
- ▶ For the empty residue,  $R_i'$  is fail
  - ▶ Then the use of the residue  $R_i$  entails the addition of  $R_i'$  to the definition of  $H$ .



# An Example Scenario

---

EDB: Ships(z1,z2,z3,z4,z5,z6,z7)

Owners(y1,y2,y3,y4)

IC:

IC1:  $\leftarrow$  Owners(y1,'iceland',y3,y4)

IC2:  $z2 = \text{'onassis'} \leftarrow$  Ships(z1,z2, 'supertanker',z4,z5,z6,z7)

IC+:

IC1+:  $\leftarrow$  Owners(y1,x1,y3,y4), x1='iceland'

IC2+:  $z2 = \text{'onassis'} \leftarrow$  Ships(z1,z2,x1,z4,z5,z6,z7), x1='supertanker'

# An Example Scenario

## ► For IC1+

← Owners(y1,x1,y3,y4), (x1='iceland')

Owners(k1,k2,k3,k4) ←

{k1/y1, k2/x1, k3/y3, k4/y4}

← k2='iceland'

## ► For IC2+

z2 = 'onassis' ← Ships(z1,z2,x1,z4,z5,z6,z7),  
x1='supertanker'

Ships(k1,k2,k3,k4,k5,k6,k7)←

{k1/z1, k2/z1, k3/z3, k4/z4, k5/z5, k6/z6, k7/z7}

k2='onassis' ← k3='superthanker'

# An Example Scenario

---

- ▶ The semantically constrained axioms (SCA) are:

SCA1:  $\text{Ships}(z1, z2, z3, z4, z5, z6, z7) \leftarrow$   
 $\text{Ships}(z1, z2, z3, z4, z5, z6, z7) \{ z2 = \text{'onassis'} \leftarrow z3 = \text{'superthanker'} \}$

SCA2:  $\text{Owners}(y1, y2, y3, y4) \leftarrow \text{Owners}(y1, y2, y3, y4) \{ \leftarrow y2 = \text{'iceland'} \}$

- ▶ The meanings of these SCA are as follows:

for SCA1:

$\text{Ships}(z1, z2, z3, z4, z5, z6, z7) \leftarrow$   
 $\text{Ships}(z1, z2, z3, z4, z5, z6, z7), \text{not}(z3 = \text{'superthanker'}, \text{not}(z2 = \text{'onassis'}))$

for SCA2:

$\text{Owners}(y1, y2, y3, y4) \leftarrow \text{Owners}(y1, y2, y3, y4), \text{not}(y2 = \text{'iceland'})$