

Deductive Databases

Data Exchange

Outline

- ▶ Introduction
 - ▶ Definitions
- ▶ Schema Mapping
 - ▶ Schema mapping specification language
 - ▶ Examples
- ▶ Data Exchange Framework
 - ▶ Universal solution of Data Exchange
 - ▶ Examples
- ▶ Weakly Acyclic Sets of Tgds
 - ▶ Position Graph

Introduction

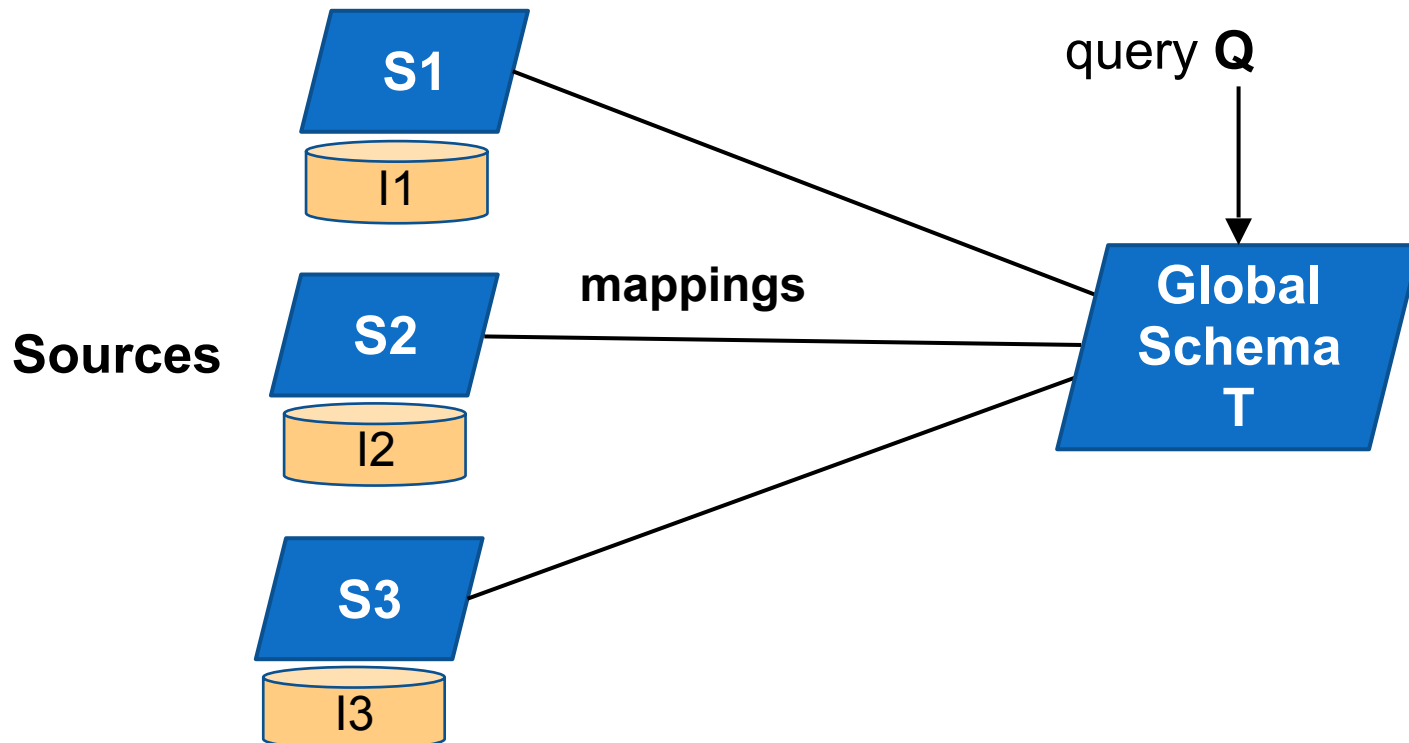
- ▶ Data is inherently heterogeneous
 - ▶ Due to the explosion of online data repositories
 - ▶ Due to the variety of users, who develop a wealth of applications
 - ▶ at different time
 - ▶ with disparate requirements in their mind

Introduction

- ▶ A fundamental requirement is to translate data across different formats and to ensure data interoperability
- ▶ Two facets of the same problem
 - ▶ Data Integration
 - ▶ Data Exchange
- ▶ Other important related problems
 - ▶ Schema Integration
 - ▶ Schema Evolution

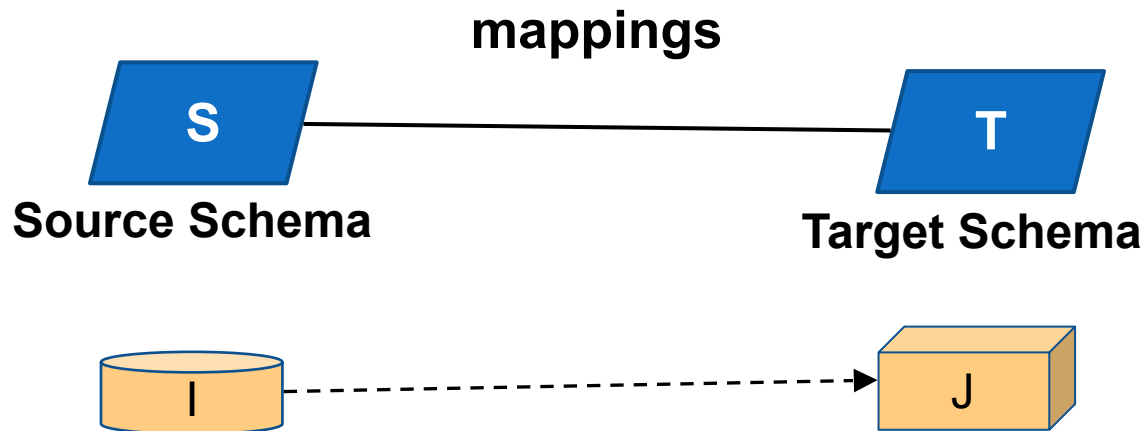
Data Integration

- ▶ Data Integration [Lenzerini 2002]
 - ▶ Query heterogeneous data in different sources via a virtual global schema



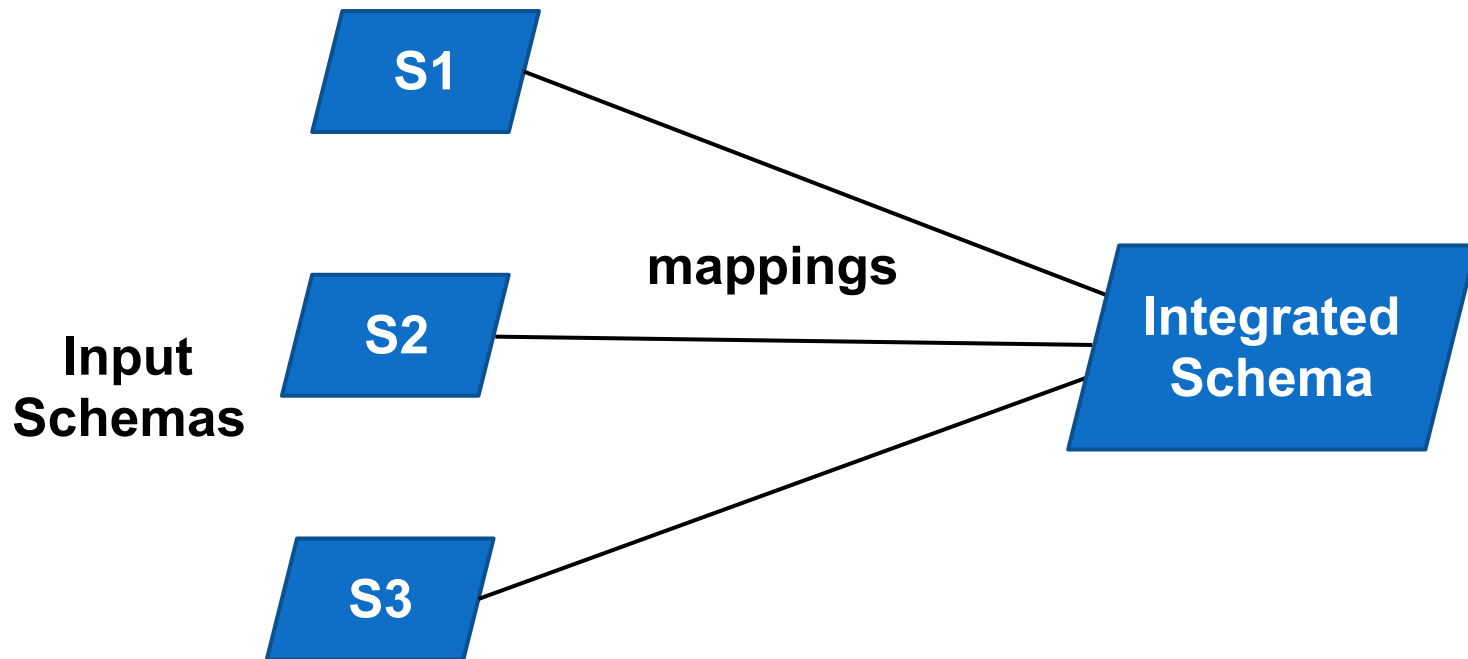
Data Exchange

- ▶ Data Exchange [Fagin et al. 2005]
 - ▶ Transform data structured under a source schema into data structured under a different target schema




Schema Integration

- ▶ Schema Integration [Batini et al. 1986]
 - ▶ A set of source schemas need to be integrated into one mediated schema



Schema Evolution

- ▶ Schema Evolution [Lerner 2000]
 - ▶ An original schema $S1$ evolves into subsequent versions $S1'$, $S1''$ etc.




Data Exchange

- ▶ Data Exchange is an old, but recurrent, database problem
- ▶ Phil Bernstein, Microsoft - 2003 “Data exchange is the oldest database problem”
 - ▶ EXPRESS: IBM San Jose Research Lab - 1977
EXtraction, Processing, and REStructuring System for transforming data between hierarchical databases.
- ▶ Data Exchange underlies: Data Warehousing, ETL (Extract-Transform-Load) tasks; XML Publishing, XML Storage; more recently, exporting relational data to RDF.

Schema Mapping

- ▶ **Schema mappings:** high-level, declarative assertions that specify the relationship between two schemas
- ▶ Ideally, schema mappings should be
 - ▶ Expressive enough to specify data interoperability tasks
 - ▶ Simple enough to be efficiently manipulated by tools
- ▶ Schema mappings constitute the essential building blocks in formalizing data integration and data exchange
- ▶ Schema mapping help with the development of tools:
 - ▶ are easier to generate and manage (semi)-automatically;
 - ▶ can be compiled into SQL/XSLT scripts automatically

Schema Mapping



- ▶ **Schema Mapping** $M = (S, T, \Sigma)$
 - ▶ Source Schema S
 - ▶ Target Schema T
 - ▶ High-level, declarative assertions Σ that specify the relationship between S and T
- ▶ **Data Exchange** via the schema mapping $M = (S, T, \Sigma)$
 - ▶ Transform a given source instance I to a target instance J , so that $\langle I, J \rangle$ satisfy the specifications Σ of M

Solutions in schema mappings

- ▶ **Definition:** Schema Mapping $M = (S, T, \Sigma)$ If I is a source instance, then a solution for I is a target instance J such that (I, J) satisfy Σ
- ▶ **Fact:** In general, for a given source instance I
 - ▶ *No* solution for I may exist or
 - ▶ *Multiple* solutions for I may exist; in fact, infinitely many solutions for I may exist

Schema Mapping specification languages

- ▶ **Question:** How are schema mappings specified?
- ▶ **Answer:** Use logic. In particular, it is natural to try to use first-order logic as a specification language for schema mappings
- ▶ **Fact:** There exists a fixed first-order sentence specifying a schema mapping M^* such that $\text{Sol}(M^*)$ is undecidable
 - ▶ Reason: undecidability of validity in FOL
 - ▶ Hence, we need to restrict ourselves to well-behaved fragments of first-order logic

Embedded Dependencies

- ▶ **Dependency Theory:** extensive study of constraints in relational databases in the 1970s and 1980s
- ▶ **Embedded Implicational Dependencies:** R. Fagin, C. Beeri and M. Vardi, ...
 - ▶ Class of constraints with a balance between high expressive power and good algorithmic properties
 - ▶ **Tuple-generating dependencies (tgds):**
 - ▶ Inclusion and multi-valued dependencies are a special case
 - ▶ **Equality-generating dependencies (egds):**
 - ▶ Functional dependencies are a special case

Schema Mapping specification languages

- ▶ How the relationship between source and target is given by formulas of first-order logic,
 - ▶ **Source-to-Target** Tuple Generating Dependencies (s-t tgds)
- ▶ $\varphi(x) \rightarrow \exists y \Psi(x, y)$
 - ▶ $\varphi(x)$ is a conjunction of atoms over the source;
 - ▶ $\Psi(x, y)$ is a conjunction of atoms over the target
- ▶ Example
 - ▶ $(\text{Student}(s) \wedge \text{Enrolls}(s, c)) \rightarrow \exists t \exists g (\text{Teachers}(t, c) \wedge \text{Grades}(s, c, g))$

Examples of simple mapping tasks

- ▶ Let us consider some simple tasks that a schema mapping specification language should support:
 - ▶ `Copy (Nicknaming)`: Copy each source table to a target table and rename it
 - ▶ `Projection`: Form a target table by projecting on one or more columns of a source table
 - ▶ `Decomposition`: Decompose a source table into two or more target tables
 - ▶ ...

Examples of simple mapping tasks

- ▶ Let us consider some simple tasks that a schema mapping specification language should support:
 - ▶ ...
 - ▶ `Column Augmentation`: Form a target table by adding one or more columns to a source table.
 - ▶ `Join`: Form a target table by joining two or more source tables
 - ▶ Combination of the above
 - ▶ e.g., “`join + column augmentation`”

Examples of simple mapping tasks

- ▶ Copy (Nicknaming) :

- ▶ Copy each source table to a target table and rename it

- ▶ $\forall x_1, \dots, x_n \left(P(x_1, \dots, x_n) \rightarrow R(x_1, \dots, x_n) \right)$

- ▶ Projection:

- ▶ Form a target table by projecting on one or more columns of a source table

- ▶ $\forall x, y, z \left(P(x, y, z) \rightarrow R(x, y) \right)$

- ▶ Decomposition:

- ▶ Decompose a source table into two or more target tables

- ▶ $\forall x, y, z \left(P(x, y, z) \rightarrow \left(R(x, y) \wedge T(y, z) \right) \right)$

Examples of simple mapping tasks

- ▶ Column Augmentation:

- ▶ Form a target table by adding one or more columns to a source table

- ▶ $\forall x, y \left(P(x, y) \rightarrow \exists z R(x, y, z) \right)$

- ▶ Join:

- ▶ Form a target table by joining two or more source tables

- ▶ $\forall x, y, z \left(E(x, z) \wedge F(z, y) \rightarrow R(x, y, z) \right)$

- ▶ join + column augmentation:


- ▶ Combination of the above

- ▶ $\forall x, y, z \left(E(x, z) \wedge F(z, y) \rightarrow \exists w T(x, y, z, w) \right)$

Target Dependencies

- ▶ In addition to source-to-target dependencies, we also consider **target dependencies**
- ▶ **Target Tgds:** $\varphi^T(x) \rightarrow \exists y \Psi^T(x, y)$
- ▶ An example of target inclusion dependency constraint
 - ▶ $\text{Dept}(did, dname, mgr_id, mgr_name) \rightarrow \text{Mgr}(mgr_id, did)$
- ▶ **Target Egds:** $\varphi^T(x) \rightarrow (x1 = x2)$
- ▶ An example of target key constraint
 - ▶ $(\text{Mgr}(e, d1) \wedge \text{Mgr}(e, d2)) \rightarrow (d1 = d2)$

Data Exchange Framework



- ▶ **Schema Mapping** $M = (S, T, \Sigma_{st}, \Sigma_t)$
 - ▶ Σ_{st} is a set of source-to-target tgds
 - ▶ Σ_t is a set of target tgds and target egds

Multiple Solutions

- ▶ **Fact:** Given a source instance, multiple solutions may exist
- ▶ **Example:**
 - ▶ Source relation $E(A,B)$
 - ▶ target relation $H(A,B)$

$$\Sigma: E(x, y) \rightarrow \exists z (H(x, z) \wedge H(z, y))$$

- ▶ **Source instance:**
 - ▶ $I = \{E(a, b)\}$
- ▶ **Solutions:**
 - ▶ Infinitely many solutions exist

Multiple Solutions: Example

- ▶ Consider a set of source-to-target dependencies Σ_{st} :
 - ▶ (d1) $\text{EmpCity}(e, c) \rightarrow \boxed{?}H \text{ Home}(e, H)$
 - ▶ (d2) $\text{EmpCity}(e, c) \rightarrow \boxed{?}D \text{ (EmpDept}(e, D) \wedge \text{DepCity}(D, c))$
 - ▶ (d3) $\text{LivesIn}(e, h) \rightarrow \text{Home}(e, h)$
 - ▶ (d4) $\text{LivesIn}(e, h) \rightarrow \exists D \exists C (\text{EmpDept}(e, D) \wedge \text{DepCity}(D, C))$
- ▶ and a source instance I such that:
 - ▶ $I = \{ \text{EmpCity}(\text{Alice}, \text{SJ})$
 $\text{EmpCity}(\text{Bob}, \text{SD})$
 $\text{LivesIn}(\text{Alice}, \text{SF})$
 $\text{LivesIn}(\text{Bob}, \text{LA}) \}$
- ▶ *Which possible solution do exist?*

Multiple Solutions: Example

A possible solution:

$J0 = \{$ Home (Alice, SF) ,
Home (Bob, LA) ,
EmpDept (Alice, D1) ,
EmpDept (Bob, D2) ,
DepCity (D1, SJ) ,
DepCity (D2, SD) }

$l:$ LivesIn (Alice, SF)
LivesIn (Bob, LA)
 $(d3):$ LivesIn (e, h) \rightarrow Home (e, h)

$l:$ EmpCity (Alice, SJ)
EmpCity (Bob, SD)
 $(d2):$ EmpCity (e, c) \rightarrow ? D (EmpDept (e, D) \wedge
DepCity (D, c))

Multiple Solutions: Example

A possible solution:

$J0' = \{$

- Home (Alice, SF) ,
- Home (Bob, LA) ,

$\{$

- EmpDept (Alice, D) ,
- EmpDept (Bob, D) ,
- DepCity (D, SJ) ,
- DepCity (D, SD) }

l: LivesIn (Alice, SF)
LivesIn (Bob, LA)
(d3): LivesIn (e, h) \rightarrow Home (e, h)

l: EmpCity (Alice, SJ)
EmpCity (Bob, SD)
(d2): EmpCity (e, c) \rightarrow

$\boxed{?} D (\text{EmpDept} (e, D) \wedge \text{DepCity} (D, c))$

Multiple Solutions: Example

A possible solution:

$J = \{ \text{Home}(\text{Alice}, \text{SF}),$
 $\text{Home}(\text{Bob}, \text{LA}),$

$\text{Home}(\text{Alice}, \text{H1}),$
 $\text{Home}(\text{Bob}, \text{H2}),$

I: $\text{LivesIn}(\text{Alice}, \text{SF})$
 $\text{LivesIn}(\text{Bob}, \text{LA})$

(d3): $\text{LivesIn}(e, h) \rightarrow \text{Home}(e, h)$

$\text{EmpDept}(\text{Alice}, \text{D1}),$
 $\text{EmpDept}(\text{Bob}, \text{D2}),$
 $\text{DepCity}(\text{D1}, \text{SJ}),$
 $\text{DepCity}(\text{D2}, \text{SD}) \}$

I: $\text{EmpCity}(\text{Alice}, \text{SJ})$
 $\text{EmpCity}(\text{Bob}, \text{SD})$

(d1): $\text{EmpCity}(e, c) \rightarrow$? $\text{Home}(e, H)$

I: $\text{EmpCity}(\text{Alice}, \text{SJ})$
 $\text{EmpCity}(\text{Bob}, \text{SD})$

(d2): $\text{EmpCity}(e, c) \rightarrow$? $\text{D}(\text{EmpDept}(e, D) \wedge$
 $\text{DepCity}(D, c))$

Main issues in Data Exchange

- ▶ For a given source instance, there may be multiple target instances satisfying the specifications of the schema mapping
- ▶ When more than one solution exist, which solutions are “better” than others?
- ▶ How do we compute a “best” solution?
- ▶ In other words, what is the “right” semantics of data exchange?

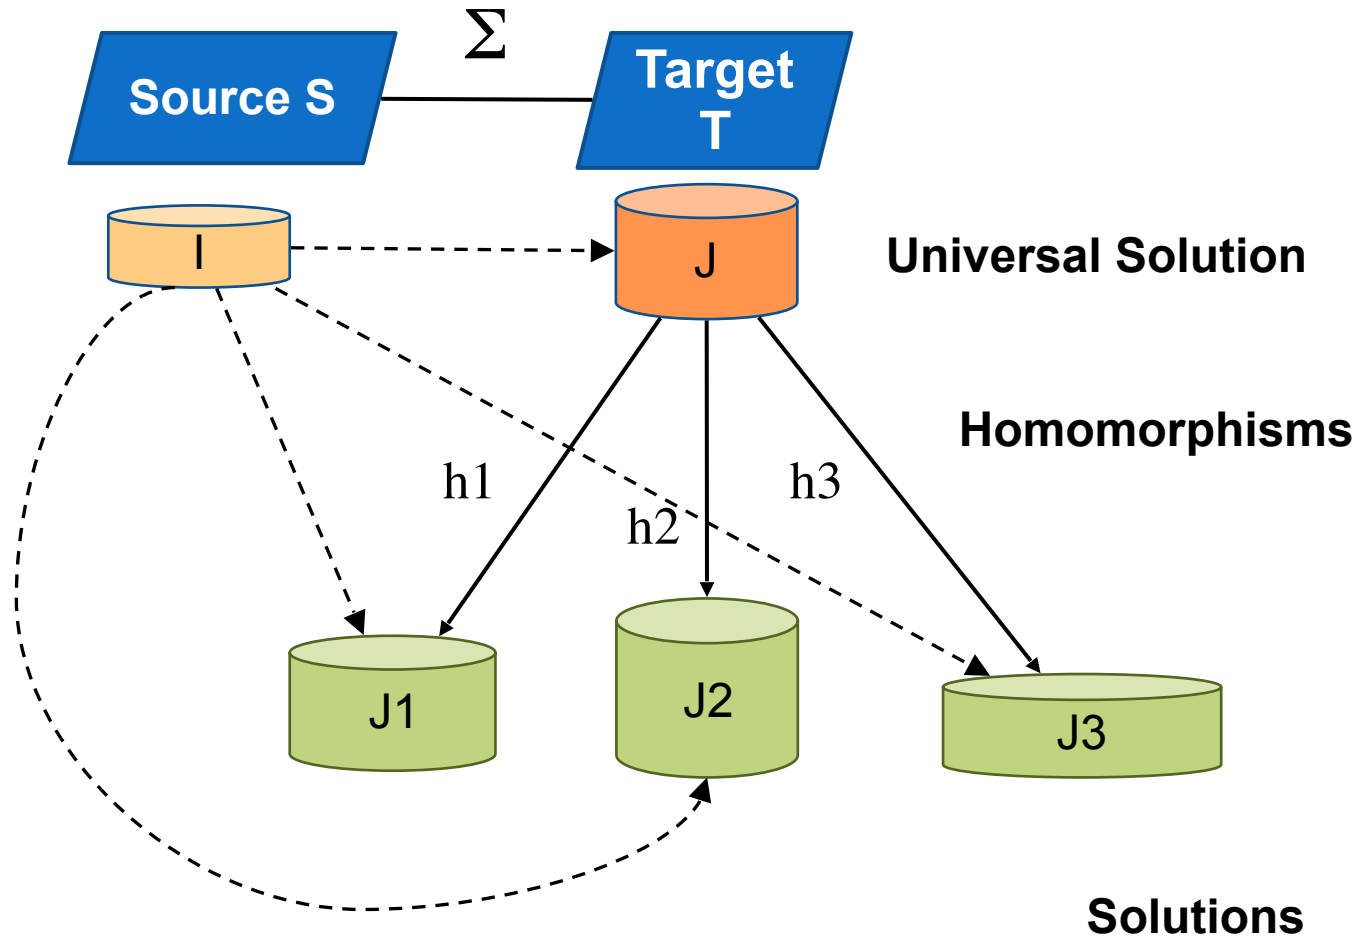
Universal solution in Data Exchange

- ▶ We introduce the notion of **universal solutions** as the “best” solutions in data exchange
- ▶ By definition, a solution is universal if it has homomorphisms to all other solutions
 - ▶ It is a “most general” solution
- ▶ Constants
 - ▶ entries in source instances
- ▶ Variables (labeled nulls)
 - ▶ invented entries in target instances

Universal solution in Data Exchange

- ▶ By definition, a solution is universal if it has homomorphisms to all other solutions
 - ▶ It is a “most general” solution
- ▶ Homomorphism $h: J1 \rightarrow J2$ between target instances:
 - ▶ $h(c) = c$, for every constant c in $J1$
 - ▶ For every fact $P(a_1, \dots, a_m)$ in $J1$, then we have that $P(h(a_1), \dots, h(a_m))$ is a fact in $J2$

Schema Mapping



Structural Properties of Universal Solutions

- ▶ Universal solutions are analogous to most general unifiers in logic programming
- ▶ Uniqueness up to homomorphic equivalence: If J and J' are universal for I , then they are homomorphically equivalent
- ▶ Representation of the entire space of solutions: Assume that J is universal for I , and J' is universal for I' . Then the following are equivalent:
 - ▶ I and I' have the same space of solutions.
 - ▶ J and J' are homomorphically equivalent.

Weakly Acyclic Sets of Tgds: Definition

- ▶ **Position graph** (dependency graph) of a set Σ of tgds:
 - ▶ **Nodes:** $R.A$, with R relation symbol, A attribute of R
 - ▶ **Edges:** for every $\varphi(x) \rightarrow \exists y \Psi(x, y)$ in Σ , for every occurrence of x in φ in position $R.A$
 - ▶ for every x in x occurring in Ψ in position $S.B$
 - add an edge $R.A \rightarrow S.B$ (if it does not already exist)
 - ▶ for every existentially quantified variable y in Ψ and for every occurrence y in Ψ in position $T.C$
 - add a special edge $R.A \rightarrow T.C$ (if it does not already exist)
- is **weakly acyclic** if the position graph has **no cycle** containing a **special edge**

Weakly Acyclic Sets of Tgds

- ▶ Weakly acyclic sets of tgds contain as special cases
 - ▶ Sets of full tgds


$$\varphi^T(x, x') \rightarrow \Psi^T(x)$$

- ▶ $\varphi^T(x, x')$ and $\Psi^T(x)$ are conjunctions of atoms
- ▶ Acyclic sets of inclusion dependencies
 - ▶ Large class of dependencies occurring in practice

Weakly Acyclic Sets of Tgds: Examples

▶ Example 1: $D(d1, d2)$ and $M(m1)$

$$\Sigma = \{D(e, \textcircled{m}) \rightarrow M(\textcircled{m}), M(\textcircled{m}) \rightarrow \exists e D(\textcircled{e}, \textcircled{m})\}$$



▶ Example 2: $E(e1, e2)$

$$\Sigma = \{E(x, \textcircled{y}) \rightarrow \exists z E(\textcircled{y}, \textcircled{z})\}$$

